

CoS Platform Test Report

This document is a test fixture for the CoS ingestion pipeline.
It contains structured content to verify PDF extraction via Apache Tika.

Section 1: Platform Overview

The Chief of Staff AI platform ingests documents of many formats.
Each document is extracted to Markdown, chunked by token count,
embedded with a vector model, and stored transactionally in Postgres.

Section 2: Ingestion Pipeline

The ingestion pipeline runs: extract to markdown, chunk, embed, and store.
Documents flow through the pipeline in a single atomic transaction.
On failure the transaction is rolled back and no partial records remain.

Section 3: Provenance Tracking

Every document version is recorded with file hash and ingested timestamp.
Re-ingesting the same source path creates a new version record.
The originals directory stores byte-for-byte copies of every source file.

Section 4: Retrieval Readiness

Once ingested, documents are available for keyword and semantic search.
Each stored chunk preserves the source document id, path, and chunk index.